

Capstone Project

Image classifier for the SVHN dataset

Instructions

In this notebook, you will create a neural network that classifies real-world images digits. You will use concepts from throughout this course in building, training, testing, validating and saving your Tensorflow classifier model.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

How to submit


When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (File -> Download as -> PDF via LaTeX). You should then submit this pdf for review.

Let's get started!

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

```
In [1]: import tensorflow as tf
from scipy.io import loadmat
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelBinarizer
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pandas as pd
import random
from PIL import Image
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Softmax, Conv2D, MaxPooling2D, BatchNormalization
from tensorflow.keras import regularizers
from tensorflow.keras.callbacks import ModelCheckpoint

%matplotlib inline
```

 SVHN overview image For the capstone project, you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images in all, and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

- Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. "Reading Digits in Natural Images with Unsupervised Feature Learning". NIPS Workshop on Deep Learning and Unsupervised Feature

Learning, 2011.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a real-world image into one of ten classes.

In [2]:

```
# Run this cell to load the dataset

train = loadmat('data/train_32x32.mat')
test = loadmat('data/test_32x32.mat')
```

Both `train` and `test` are dictionaries with keys `X` and `y` for the input images and labels respectively.

1. Inspect and preprocess the dataset

- Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.
- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.
- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. *Hint: retain the channel dimension, which will now have size 1.*
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

In [3]:

```
# Lets check out the data set first

print(train.keys())
print(train['__header__'])
print(train['__version__'])
print(train['__globals__'])
print("\n***** X\n", train['X'])
print("\n***** Y\n", train['y'])
```

```
dict_keys(['__header__', '__version__', '__globals__', 'X', 'y'])
b'MATLAB 5.0 MAT-file, Platform: GLNXA64, Created on: Mon Dec 5 21:09:26 2011'
1.0
[]
```

```
***** X
[[[ 33  84  19 ...  92 190 216]
 [ 30  76  54 ...  78 188 217]
 [ 38  59 110 ... 101 191 212]]

[[ 15  86  20 ...  94 205 221]
 [ 23  73  52 ...  82 203 222]
 [ 19  66 111 ... 105 206 217]]

[[ 15  77  25 ... 114 220 226]
 [ 17  78  57 ... 101 218 227]
 [ 19  56 116 ... 125 220 221]]

...

[[ 72  90  65 ... 200 229 200]
 [ 65  78 144 ... 201 231 199]
 [ 56  69 223 ... 203 224 191]]

[[ 82  88  78 ... 192 229 193]
 [ 77  77 148 ... 193 229 188]
 [ 57  67 218 ... 195 224 182]]
```

```
[ [ 89  88  98 ... 190 229 197]
  [ 79  78 158 ... 191 228 189]
  [ 59  66 220 ... 193 223 186]]]
```

```
[[ [ 28  85  21 ...  92 183 204]
   [ 39  77  53 ...  78 182 205]
   [ 35  61 110 ... 103 186 202]]]
```

```
[ [ 14  83  19 ...  93 200 210]
  [ 25  73  52 ...  80 199 211]
  [ 22  64 106 ... 106 201 208]]]
```

```
[ [ 14  74  25 ... 111 218 220]
  [ 20  69  56 ...  98 217 221]
  [ 17  59 111 ... 124 218 217]]]
```

...

```
[ [ 40  89  63 ... 181 227 201]
  [ 39  82 137 ... 180 228 199]
  [ 50  64 208 ... 184 223 193]]]
```

```
[ [ 67  88  91 ... 177 227 195]
  [ 58  79 153 ... 176 226 191]
  [ 52  70 214 ... 180 222 186]]]
```

```
[ [ 83  88 130 ... 183 228 196]
  [ 78  81 180 ... 182 224 190]
  [ 60  67 229 ... 187 221 186]]]
```

```
[[ [ 40  83  21 ...  99 171 198]
   [ 41  76  53 ...  84 170 198]
   [ 38  60 110 ... 112 175 197]]]
```

```
[ [ 18  78  20 ...  94 189 202]
  [ 21  77  51 ...  81 189 202]
  [ 26  58 106 ... 110 193 201]]]
```

```
[ [ 16  61  22 ... 107 213 212]
  [ 17  50  52 ...  94 213 211]
  [ 23  54 106 ... 123 215 210]]]
```

...

```
[ [ 23  90  79 ... 167 231 203]
  [ 29  85 147 ... 166 230 200]
  [ 45  63 210 ... 171 226 196]]]
```

```
[ [ 35  88 125 ... 172 229 198]
  [ 42  83 181 ... 171 226 194]
  [ 44  66 230 ... 176 223 191]]]
```

```
[ [ 72  85 178 ... 185 227 195]
  [ 69  82 218 ... 184 223 190]
  [ 53  70 254 ... 189 220 187]]]
```

...

```
[[ [ 86 100  88 ...  99 187 233]
   [ 81  98 162 ...  94 185 226]
   [ 75  72 237 ... 110 186 228]]]
```

```

[[ 87  98  89 ...  96 204 230]
 [ 82  94 163 ...  91 202 224]
 [ 71  76 238 ... 109 199 225]]

[[ 82  95  84 ... 108 217 228]
 [ 79  93 156 ... 103 217 223]
 [ 65  73 230 ... 124 210 221]]

...

[[104 104  62 ... 210 204 198]
 [104 104 142 ... 207 200 196]
 [ 87  86 227 ... 204 195 190]]

[[104 102  67 ... 206 196 184]
 [105 102 144 ... 202 193 183]
 [ 81  87 226 ... 200 189 177]]

[[103 100  74 ... 203 196 189]
 [105 101 145 ... 197 193 187]
 [ 78  78 225 ... 199 189 182]]]

[[[ 84 103  88 ...  94 186 231]
 [ 86 104 164 ...  91 184 226]
 [ 64  79 240 ... 103 185 228]]

[[ 86 106  87 ...  94 198 229]
 [ 79 104 160 ...  91 197 224]
 [ 72  79 237 ... 104 194 225]]

[[ 82 103  88 ... 110 211 227]
 [ 76 103 159 ... 107 211 223]
 [ 72  87 237 ... 121 204 222]]

...

[[110 103  60 ... 219 222 195]
 [103 104 141 ... 218 216 194]
 [ 84  86 230 ... 215 212 186]]

[[106 103  61 ... 218 214 181]
 [105 103 141 ... 215 209 181]
 [ 85  87 228 ... 212 205 173]]

[[106 105  65 ... 212 208 186]
 [104  99 143 ... 209 205 183]
 [ 86  81 226 ... 209 200 177]]]

[[[ 85 103  84 ...  88 190 230]
 [ 88 106 160 ...  87 188 226]
 [ 68  82 238 ...  94 190 227]]

[[ 89 103  81 ...  85 199 230]
 [ 82 105 154 ...  84 197 226]
 [ 72  87 233 ...  93 194 227]]

[[ 85 104  87 ... 105 208 229]
 [ 79 106 158 ... 103 208 225]
 [ 67  91 238 ... 114 201 226]]

...

[[111 113  63 ... 217 232 190]

```

```

[104 103 144 ... 217 227 190]
[ 87  88 235 ... 214 223 181]]

[[109 104  62 ... 221 226 178]
 [105 104 143 ... 220 221 177]
 [ 86  88 232 ... 219 216 169]]

[[103 103  63 ... 218 218 181]
 [106  98 145 ... 217 213 178]
 [ 79  80 231 ... 218 209 171]]]]

***** Y
[[1]
 [9]
 [2]
 ...
 [1]
 [6]
 [9]]

```

In [4]: *#Let's not extract the images and the labels from both the train and test dictionaries and*

```

train_x = train['X']
train_y = train['y']
test_x = test['X']
test_y = test['y']

print("Shape of Training Data: ", train_x.shape)
print("Shape of Test Data: ", train_x.shape)
print("Shape of Training Lables: ", train_y.shape)
print("Shape of Test Lables: ", test_y.shape)

print("\nData at train_y[1]:",train_y[1])
print("\nData at train_x[1]:\n",train_x[1])

```

```

Shape of Training Data: (32, 32, 3, 73257)
Shape of Test Data: (32, 32, 3, 73257)
Shape of Training Lables: (73257, 1)
Shape of Test Lables: (26032, 1)

```

```
Data at train_y[1]: [9]
```

```

Data at train_x[1]:
[[[ 28  85  21 ...  92 183 204]
 [ 39  77  53 ...  78 182 205]
 [ 35  61 110 ... 103 186 202]]

[[ 14  83  19 ...  93 200 210]
 [ 25  73  52 ...  80 199 211]
 [ 22  64 106 ... 106 201 208]]

[[ 14  74  25 ... 111 218 220]
 [ 20  69  56 ...  98 217 221]
 [ 17  59 111 ... 124 218 217]]

...

[[ 40  89  63 ... 181 227 201]
 [ 39  82 137 ... 180 228 199]
 [ 50  64 208 ... 184 223 193]]

[[ 67  88  91 ... 177 227 195]
 [ 58  79 153 ... 176 226 191]
 [ 52  70 214 ... 180 222 186]]

```

```
[[ 83  88 130 ... 183 228 196]
 [ 78  81 180 ... 182 224 190]
 [ 60  67 229 ... 187 221 186]]]
```

```
In [5]: print("Unique Training Labels: ", np.unique(train_y))
        print("Unique Test Labels: ", np.unique(test_y))
```

```
Unique Training Labels: [ 1  2  3  4  5  6  7  8  9 10]
Unique Test Labels:    [ 1  2  3  4  5  6  7  8  9 10]
```

```
In [6]: train_y[train_y == 10] = 0
        test_y[test_y == 10] = 0
        print("Unique Training Labels: ", np.unique(train_y))
        print("Unique Test Labels: ", np.unique(test_y))
        print("train_y [2]: " , train_y[2600])
        print("train_y [2]: " , train_y[600])
```

```
Unique Training Labels: [0 1 2 3 4 5 6 7 8 9]
Unique Test Labels:    [0 1 2 3 4 5 6 7 8 9]
train_y [2]: [1]
train_y [2]: [7]
```

```
In [7]: train_x_std = train_x/255.
        test_x_std = test_x/255.
```

```
In [8]: train_x_std = np.moveaxis(train_x_std, -1, 0)
        test_x_std = np.moveaxis(test_x_std, -1, 0)
        print("Shape of Training Data: ", train_x_std.shape)
        print("Shape of Test Data: ", test_x_std.shape)
        print("Shape of Training Lables: ", train_y.shape)
        print("Shape of Test Lables: ", test_y.shape)
```

```
Shape of Training Data: (73257, 32, 32, 3)
Shape of Test Data:    (26032, 32, 32, 3)
Shape of Training Lables: (73257, 1)
Shape of Test Lables:   (26032, 1)
```

```
In [9]: train_x_std
```

```
Out[9]: array([[[[0.12941176, 0.11764706, 0.14901961],
                  [0.05882353, 0.09019608, 0.0745098 ],
                  [0.05882353, 0.06666667, 0.0745098 ],
                  ...,
                  [0.28235294, 0.25490196, 0.21960784],
                  [0.32156863, 0.30196078, 0.22352941],
                  [0.34901961, 0.30980392, 0.23137255]],
                  [0.10980392, 0.15294118, 0.1372549 ],
                  [0.05490196, 0.09803922, 0.08627451],
                  [0.05490196, 0.07843137, 0.06666667],
                  ...,
                  [0.15686275, 0.15294118, 0.19607843],
                  [0.2627451 , 0.22745098, 0.20392157],
                  [0.3254902 , 0.30588235, 0.23529412]],
                  [0.15686275, 0.16078431, 0.14901961],
                  [0.07058824, 0.08235294, 0.10196078],
                  [0.0627451 , 0.06666667, 0.09019608],
                  ...,
                  [0.09019608, 0.11372549, 0.17647059],
                  [0.1372549 , 0.16470588, 0.17254902],
                  [0.28235294, 0.27058824, 0.20784314]]],
```

```

...,

[[0.3372549 , 0.31764706, 0.29411765],
 [0.34117647, 0.32156863, 0.27843137],
 [0.32156863, 0.30980392, 0.25490196],
 ...,
 [0.40784314, 0.40784314, 0.34117647],
 [0.40784314, 0.41176471, 0.31764706],
 [0.40392157, 0.41176471, 0.30588235]],

[[0.32941176, 0.3372549 , 0.25098039],
 [0.3372549 , 0.30980392, 0.28235294],
 [0.32156863, 0.29803922, 0.28235294],
 ...,
 [0.43137255, 0.40392157, 0.32941176],
 [0.41568627, 0.41176471, 0.33333333],
 [0.41568627, 0.40784314, 0.3372549 ]],

[[0.33333333, 0.34509804, 0.26666667],
 [0.34901961, 0.32156863, 0.28235294],
 [0.33333333, 0.30980392, 0.2627451 ],
 ...,
 [0.43529412, 0.40784314, 0.34117647],
 [0.42745098, 0.41176471, 0.3372549 ],
 [0.40392157, 0.41568627, 0.30980392]]],

[[[0.32941176, 0.29803922, 0.23137255],
 [0.3372549 , 0.28627451, 0.25882353],
 [0.30196078, 0.30588235, 0.21960784],
 ...,
 [0.35294118, 0.30588235, 0.27058824],
 [0.34509804, 0.30196078, 0.2627451 ],
 [0.34509804, 0.30588235, 0.25882353]],

[[0.33333333, 0.30196078, 0.23921569],
 [0.3254902 , 0.28627451, 0.25098039],
 [0.29019608, 0.27058824, 0.23137255],
 ...,
 [0.34901961, 0.32156863, 0.25098039],
 [0.34509804, 0.30980392, 0.2745098 ],
 [0.34509804, 0.31764706, 0.2627451 ]],

[[0.3254902 , 0.29803922, 0.23529412],
 [0.30588235, 0.30196078, 0.22745098],
 [0.23921569, 0.19607843, 0.21176471],
 ...,
 [0.35294118, 0.33333333, 0.24705882],
 [0.34509804, 0.3254902 , 0.25882353],
 [0.33333333, 0.32156863, 0.2745098 ]],

...,

[[0.39215686, 0.38431373, 0.28235294],
 [0.38431373, 0.36862745, 0.29803922],
 [0.37254902, 0.36470588, 0.28627451],
 ...,
 [0.40784314, 0.40784314, 0.3372549 ],
 [0.4 , 0.4 , 0.34117647],
 [0.39215686, 0.39607843, 0.30588235]],

[[0.40392157, 0.40784314, 0.30980392],
 [0.41568627, 0.40784314, 0.30980392],
 [0.40392157, 0.40392157, 0.34117647],
 ...,

```

```
[0.40392157, 0.40784314, 0.3372549 ],
[0.40392157, 0.40392157, 0.34117647],
[0.41176471, 0.38823529, 0.31764706]],

[[[0.40392157, 0.41568627, 0.32156863],
[0.40392157, 0.41176471, 0.34117647],
[0.40784314, 0.41568627, 0.35686275],
...,
[0.44313725, 0.40392157, 0.34509804],
[0.40784314, 0.40784314, 0.34509804],
[0.40392157, 0.38431373, 0.31372549]]],

[[[0.0745098 , 0.21176471, 0.43137255],
[0.07843137, 0.20392157, 0.43529412],
[0.09803922, 0.22352941, 0.45490196],
...,
[0.25490196, 0.56470588, 0.8745098 ],
[0.30588235, 0.58039216, 0.85490196],
[0.38431373, 0.61960784, 0.8627451 ]],

[[[0.08235294, 0.20784314, 0.43137255],
[0.0745098 , 0.20392157, 0.41568627],
[0.09803922, 0.21960784, 0.43529412],
...,
[0.24705882, 0.5372549 , 0.81568627],
[0.35686275, 0.6 , 0.83921569],
[0.50980392, 0.70588235, 0.89803922]]],

[[[0.08235294, 0.20784314, 0.43137255],
[0.07843137, 0.2 , 0.41568627],
[0.08627451, 0.20392157, 0.41568627],
...,
[0.30980392, 0.57647059, 0.82352941],
[0.49019608, 0.70980392, 0.90196078],
[0.69803922, 0.85490196, 0.99607843]]],

...,

[[[0.34509804, 0.63529412, 0.92941176],
[0.34901961, 0.63921569, 0.93333333],
[0.32941176, 0.61176471, 0.90196078],
...,
[0.24313725, 0.55686275, 0.89019608],
[0.2627451 , 0.56470588, 0.88627451],
[0.29019608, 0.56862745, 0.88235294]]],

[[[0.34509804, 0.64313725, 0.94117647],
[0.34117647, 0.62745098, 0.92941176],
[0.34509804, 0.62352941, 0.92941176],
...,
[0.23529412, 0.55294118, 0.90196078],
[0.23921569, 0.55294118, 0.89411765],
[0.25490196, 0.56078431, 0.88627451]]],

[[[0.32941176, 0.62745098, 0.93333333],
[0.31764706, 0.60392157, 0.91372549],
[0.34117647, 0.61960784, 0.93333333],
...,
[0.24705882, 0.56470588, 0.92156863],
[0.24313725, 0.56078431, 0.90980392],
[0.24705882, 0.56862745, 0.90588235]]],

...,
```



```
[[[0.36078431, 0.30588235, 0.39607843],
  [0.36862745, 0.32156863, 0.41176471],
  [0.44705882, 0.39607843, 0.49019608],
  ...,
  [0.78431373, 0.78823529, 0.79607843],
  [0.75294118, 0.75686275, 0.76470588],
  [0.74509804, 0.74901961, 0.75686275]]],

[[[0.36078431, 0.30588235, 0.40392157],
  [0.36470588, 0.31372549, 0.41568627],
  [0.43529412, 0.38431373, 0.48627451],
  ...,
  [0.70980392, 0.70588235, 0.72156863],
  [0.69411765, 0.69019608, 0.70588235],
  [0.71764706, 0.71372549, 0.73333333]]],

[[[0.38823529, 0.32941176, 0.43921569],
  [0.36862745, 0.31764706, 0.43137255],
  [0.41960784, 0.36862745, 0.48235294],
  ...,
  [0.65490196, 0.65098039, 0.67058824],
  [0.6745098 , 0.67058824, 0.69019608],
  [0.7254902 , 0.72156863, 0.74117647]]],

...,

[[[0.38823529, 0.36862745, 0.43137255],
  [0.37647059, 0.35686275, 0.42745098],
  [0.42352941, 0.40392157, 0.48627451],
  ...,
  [0.82352941, 0.81176471, 0.8          ],
  [0.80784314, 0.79215686, 0.78431373],
  [0.79607843, 0.77254902, 0.78039216]]],

[[[0.36862745, 0.35686275, 0.40392157],
  [0.36862745, 0.35686275, 0.40784314],
  [0.43137255, 0.41960784, 0.4745098 ],
  ...,
  [0.85882353, 0.85490196, 0.84313725],
  [0.85490196, 0.84313725, 0.83137255],
  [0.83137255, 0.81960784, 0.81960784]]],

[[[0.34509804, 0.34117647, 0.36862745],
  [0.33333333, 0.32941176, 0.36470588],
  [0.41176471, 0.40392157, 0.44705882],
  ...,
  [0.85098039, 0.85098039, 0.83921569],
  [0.86666667, 0.8627451 , 0.85882353],
  [0.85490196, 0.85098039, 0.85490196]]],

[[[0.74509804, 0.7372549 , 0.74901961],
  [0.80392157, 0.79607843, 0.80784314],
  [0.8627451 , 0.85490196, 0.8627451 ],
  ...,
  [0.89803922, 0.90588235, 0.87843137],
  [0.89803922, 0.89803922, 0.87843137],
  [0.89803922, 0.89411765, 0.8745098 ]],

[[[0.71764706, 0.71372549, 0.72941176],
  [0.78431373, 0.78039216, 0.78823529],
  [0.85490196, 0.85098039, 0.85490196],
  ...,
  [0.89019608, 0.89411765, 0.8745098 ],
  [0.89019608, 0.88627451, 0.87058824],
```

```

[0.89411765, 0.87843137, 0.86666667]],

[[0.67058824, 0.66666667, 0.68627451],
[0.74117647, 0.74117647, 0.75686275],
[0.83529412, 0.83529412, 0.84313725],
...,
[0.90588235, 0.90196078, 0.88627451],
[0.89803922, 0.88627451, 0.8745098 ],
[0.89019608, 0.8745098 , 0.8627451 ]],

...,

[[0.73333333, 0.7254902 , 0.72941176],
[0.8 , 0.79215686, 0.78039216],
[0.85098039, 0.85098039, 0.82352941],
...,
[0.8 , 0.78431373, 0.76470588],
[0.76862745, 0.75686275, 0.74117647],
[0.76862745, 0.75686275, 0.74117647]],

[[0.72941176, 0.72156863, 0.7254902 ],
[0.77647059, 0.77254902, 0.76078431],
[0.82745098, 0.82745098, 0.8 ],
...,
[0.87058824, 0.84705882, 0.83137255],
[0.83921569, 0.81960784, 0.80392157],
[0.81568627, 0.80392157, 0.78431373]],

[[0.74509804, 0.7372549 , 0.74509804],
[0.78039216, 0.77254902, 0.76078431],
[0.81568627, 0.81568627, 0.78823529],
...,
[0.90980392, 0.89019608, 0.8745098 ],
[0.88627451, 0.86666667, 0.84705882],
[0.85490196, 0.83529412, 0.81960784]]],

[[[0.84705882, 0.85098039, 0.83137255],
[0.86666667, 0.87058824, 0.85098039],
[0.88627451, 0.89019608, 0.86666667],
...,
[0.78431373, 0.78039216, 0.74901961],
[0.75686275, 0.7372549 , 0.71372549],
[0.77254902, 0.74117647, 0.72941176]],

[[0.8 , 0.80392157, 0.79215686],
[0.82352941, 0.82745098, 0.81568627],
[0.8627451 , 0.86666667, 0.85098039],
...,
[0.78823529, 0.78039216, 0.75686275],
[0.76470588, 0.74901961, 0.72941176],
[0.76862745, 0.74509804, 0.72941176]],

[[0.77647059, 0.77647059, 0.77254902],
[0.79215686, 0.79215686, 0.78823529],
[0.83137255, 0.82745098, 0.82352941],
...,
[0.79607843, 0.78431373, 0.76862745],
[0.77647059, 0.76078431, 0.74901961],
[0.76470588, 0.74509804, 0.73333333]],

...,

[[0.91372549, 0.88627451, 0.89411765],
[0.90196078, 0.87843137, 0.88235294],
[0.89411765, 0.8745098 , 0.86666667],

```

```

...
[0.77647059, 0.76862745, 0.74509804],
[0.72156863, 0.71764706, 0.69411765],
[0.74117647, 0.73333333, 0.71372549]],

[[0.90588235, 0.88627451, 0.89411765],
[0.89803922, 0.87843137, 0.88235294],
[0.89019608, 0.8745098 , 0.87058824],
...
[0.76470588, 0.76078431, 0.72941176],
[0.70980392, 0.70980392, 0.67843137],
[0.72941176, 0.71764706, 0.69411765]],

[[0.90196078, 0.88627451, 0.89019608],
[0.90196078, 0.88627451, 0.89019608],
[0.89803922, 0.88235294, 0.88627451],
...
[0.74509804, 0.74509804, 0.70980392],
[0.69803922, 0.69411765, 0.6627451 ],
[0.70980392, 0.69803922, 0.67058824]]]])

```

```

In [10]: train_x_std = np.mean(train_x_std,keepdims=True,axis = -1)
test_x_std = np.mean(test_x_std,keepdims=True,axis = -1)
print("Shape of Grey Training Data: ", train_x_std.shape)

```

Shape of Grey Training Data: (73257, 32, 32, 1)

```

In [11]: train_y[train_y == 10] = 0
test_y[test_y == 10] = 0

```

```

In [12]: print("Unique Training Labels: ", np.unique(train_y))
print("Unique Test Labels: ", np.unique(test_y))

```

Unique Training Labels: [0 1 2 3 4 5 6 7 8 9]
Unique Test Labels: [0 1 2 3 4 5 6 7 8 9]

```

In [13]: lb = LabelBinarizer()

train_y = lb.fit_transform(train_y)
test_y = lb.fit_transform(test_y)

```

```

In [16]: print("Unique Training Labels: ", np.unique(train_y))
print("Unique Test Labels: ", np.unique(test_y))

```

Unique Training Labels: [0 1]
Unique Test Labels: [0 1]

```

In [17]: non_encoded_test_y = lb.inverse_transform(test_y)
non_encoded_train_y = lb.inverse_transform(train_y)

```

```

In [18]: print("Unique Training Labels: ", np.unique(non_encoded_train_y))
print("Unique Test Labels: ", np.unique(non_encoded_test_y))

```

Unique Training Labels: [0 1 2 3 4 5 6 7 8 9]
Unique Test Labels: [0 1 2 3 4 5 6 7 8 9]

```


In [19]: #Lets examine the random 10 images

for i in range(0,10):


```

```
figure, axe = plt.subplots(ncols = 1, nrows =1, figsize=(10, 1))
random_number = np.random.choice(test_x_std.shape[0])
axe.imshow(train_x_std[random_number])
axe.set_title('Actual Label: {}'.format(non_encoded_train_y[random_number]))
```


Actual Label: 1




Actual Label: 7




Actual Label: 1




Actual Label: 7




Actual Label: 2




Actual Label: 4




Actual Label: 4




Actual Label: 5



Actual Label: 3





```
In [20]: print (train_x_std.ndim, train_y.ndim, train_x_std[0].shape, train_y[0].shape)

4 2 (32, 32, 1) (10,)
```

2. MLP neural network classifier

- Build an MLP classifier model using the Sequential API. Your model should use only Flatten and Dense layers, with the final layer having a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different MLP architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 4 or 5 layers.*
- Print out the model summary (using the summary() method)
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- As a guide, you should aim to achieve a final categorical cross entropy training loss of less than 1.0 (the validation loss might be higher).
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

Regularized MLP model with Callbacks

```
In [21]: #checkpoint_path = 'model_checkpoints/checkpoint_{epoch:02d}'
! rm -r model_checkpoints
checkpoint_path = 'model_checkpoints/checkpoint'

checkpoint_val_per_epoch = ModelCheckpoint(filepath = checkpoint_path,
                                           frequency = 'epoch',
                                           monitor = 'val_loss',
                                           save_best_only = True,
                                           save_weights_only = True,
                                           verbose = 1)
```

```
In [22]: def get_mlp_regularized_model(wd, input_shape):

    model = Sequential([
        Flatten(input_shape = input_shape),
        Dense(128, kernel_regularizer = regularizers.l2(wd), activation = 'relu', name = 'layer1'),
        Dense(128, kernel_regularizer = regularizers.l2(wd), activation = 'relu', name = 'layer2'),
        Dense(128, kernel_regularizer = regularizers.l2(wd), activation = 'relu', name = 'layer3'),
        Dense(10, activation = 'softmax'),
    ])
    return model
```

```
In [23]: MLP_reg_model = get_mlp_regularized_model(1e-5, train_x_std[0].shape)
```

```

MLP_reg_model.summary()

MLP_reg_model.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics=['acc

reduceLRonPlateau_callback = tf.keras.callbacks.ReduceLRonPlateau(monitor = "val_loss",
                                                                    factor = 0.2, patience =
                                                                    verbose = 1)

MLP_reg_history = MLP_reg_model.fit(x=train_x_std, y=train_y, epochs = 30,
                                    validation_split=0.15, batch_size=64,
                                    callbacks = [checkpoint_val_per_epoch, reduceLRonPlateau_callback],
                                    verbose = False )

```

2022-05-22 17:32:40.916370: I tensorflow/core/platform/cpu_feature_guard.cc:145] This TensorFlow binary is optimized with Intel(R) MKL-DNN to use the following CPU instructions in performance critical operations: SSE4.1 SSE4.2 AVX AVX2 FMA
To enable them in non-MKL-DNN operations, rebuild TensorFlow with the appropriate compiler flags.

2022-05-22 17:32:40.916756: I tensorflow/core/common_runtime/process_util.cc:115] Creating new thread pool with default inter op setting: 8. Tune using inter_op_parallelism_threads for best performance.

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 1024)	0

layer_1 (Dense)	(None, 128)	131200

layer_2 (Dense)	(None, 128)	16512

layer_3 (Dense)	(None, 128)	16512

dense (Dense)	(None, 10)	1290
=====		
Total params: 165,514		
Trainable params: 165,514		
Non-trainable params: 0		

Epoch 00001: val_loss improved from inf to 1.48287, saving model to model_checkpoints/checkpoint

Epoch 00002: val_loss improved from 1.48287 to 1.20886, saving model to model_checkpoints/checkpoint

Epoch 00003: val_loss improved from 1.20886 to 1.17107, saving model to model_checkpoints/checkpoint

Epoch 00004: val_loss improved from 1.17107 to 1.10773, saving model to model_checkpoints/checkpoint

Epoch 00005: val_loss improved from 1.10773 to 1.06562, saving model to model_checkpoints/checkpoint

Epoch 00006: val_loss improved from 1.06562 to 1.01741, saving model to model_checkpoints/checkpoint

Epoch 00007: val_loss improved from 1.01741 to 0.97146, saving model to model_checkpoints/checkpoint

Epoch 00008: val_loss improved from 0.97146 to 0.96504, saving model to model_checkpoints/checkpoint

Epoch 00009: val_loss improved from 0.96504 to 0.92797, saving model to model_checkpoints/checkpoint

Epoch 00010: val_loss did not improve from 0.92797

Epoch 00011: val_loss did not improve from 0.92797

Epoch 00012: val_loss did not improve from 0.92797

Epoch 00012: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.

Epoch 00013: val_loss improved from 0.92797 to 0.84538, saving model to model_checkpoints/checkpoint

Epoch 00014: val_loss improved from 0.84538 to 0.82705, saving model to model_checkpoints/checkpoint

Epoch 00015: val_loss did not improve from 0.82705

Epoch 00016: val_loss improved from 0.82705 to 0.82499, saving model to model_checkpoints/checkpoint

Epoch 00017: val_loss improved from 0.82499 to 0.82358, saving model to model_checkpoints/checkpoint

Epoch 00018: val_loss improved from 0.82358 to 0.81857, saving model to model_checkpoints/checkpoint

Epoch 00019: val_loss improved from 0.81857 to 0.81474, saving model to model_checkpoints/checkpoint

Epoch 00020: val_loss improved from 0.81474 to 0.80494, saving model to model_checkpoints/checkpoint

Epoch 00021: val_loss did not improve from 0.80494

Epoch 00022: val_loss did not improve from 0.80494

Epoch 00023: val_loss did not improve from 0.80494

Epoch 00023: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.

Epoch 00024: val_loss improved from 0.80494 to 0.78978, saving model to model_checkpoints/checkpoint

Epoch 00025: val_loss improved from 0.78978 to 0.78802, saving model to model_checkpoints/checkpoint

Epoch 00026: val_loss did not improve from 0.78802

Epoch 00027: val_loss did not improve from 0.78802

Epoch 00028: val_loss did not improve from 0.78802

Epoch 00028: ReduceLROnPlateau reducing learning rate to 8.000000525498762e-06.

Epoch 00029: val_loss improved from 0.78802 to 0.78468, saving model to model_checkpoints/checkpoint

Epoch 00030: val_loss improved from 0.78468 to 0.78427, saving model to model_checkpoints/checkpoint

In [24]:

```
! ls -lh model_checkpoints/
```

total 3904


```
-rw-r--r-- 1 moni staff 77B May 22 17:36 checkpoint
-rw-r--r-- 1 moni staff 1.9M May 22 17:36 checkpoint.data-00000-of-00001
-rw-r--r-- 1 moni staff 2.0K May 22 17:36 checkpoint.index
```

```
In [25]: MLP_reg_model.evaluate(test_x_std, test_y, verbose = 2)
```

```
26032/1 - 2s - loss: 0.9941 - accuracy: 0.7295
```


```
Out[25]: [0.884822539266224, 0.7295252]
```

```
In [26]: plt.plot(MLP_reg_history.history['loss'])
plt.plot(MLP_reg_history.history['val_loss'])
plt.title('Loss vs. epochs')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')
plt.show()
```



```
In [27]: plt.plot(MLP_reg_history.history['accuracy'])
plt.plot(MLP_reg_history.history['val_accuracy'])
plt.title('Accuracy vs Epochs')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(['Training', 'Validation'], loc = 'lower right')
```

```
Out[27]: <matplotlib.legend.Legend at 0x7fb67eb45c50>
```




```
In [28]:
```




```
plt.imshow(test_x_std[3])
prediction = MLP_reg_model.predict(test_x_std[np.newaxis, 3])
print(np.argmax(prediction))
prediction
```

Out[28]: 0
array([[0.31552488, 0.11892834, 0.03276519, 0.22020902, 0.08404231,
0.08155134, 0.03576181, 0.00647927, 0.0232967 , 0.0814411]],
dtype=float32)




In [29]: plt.imshow(test_x_std[769])
prediction = MLP_reg_model.predict(test_x_std[np.newaxis, 769])
print(np.argmax(prediction))
prediction

Out[29]: 3
array([[1.05067345e-04, 7.70687219e-03, 1.12163741e-02, 7.79664874e-01,
3.74978408e-03, 1.45806253e-01, 5.92956145e-04, 5.81146684e-04,
2.01668944e-02, 3.04097552e-02]], dtype=float32)



In [30]: plt.imshow(test_x_std[2349])
prediction = MLP_reg_model.predict(test_x_std[np.newaxis, 2349])
print("Prediction: ", np.argmax(prediction))
prediction

Out[30]: Prediction: 0
array([[9.6230423e-01, 1.7757786e-03, 2.0593952e-03, 1.2357580e-03,
2.0822929e-03, 1.0418070e-03, 1.1581531e-02, 6.5745774e-04,
6.2280162e-03, 1.1033708e-02]], dtype=float32)



3. CNN neural network classifier

- Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, BatchNormalization, Flatten, Dense and Dropout layers. The final layer should again have a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different CNN architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 2 or 3 convolutional layers and 2 fully connected layers.*
- The CNN model should use fewer trainable parameters than your MLP model.
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- You should aim to beat the MLP model performance with fewer parameters!
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
In [31]: #checkpoint_path = 'model_checkpoints/checkpoint_{epoch:02d}'
! rm -r model_cnn_reg_checkpoints
checkpoint_cnn_reg_path = 'model_reg_checkpoints/checkpoint'

checkpoint_val_reg_per_epoch = ModelCheckpoint(filepath = checkpoint_cnn_reg_path,
                                              frequency = 'epoch',
                                              monitor = 'val_loss',
                                              save_best_only = True,
                                              save_weights_only = True,
                                              verbose = 1)
```

rm: model_cnn_reg_checkpoints: No such file or directory

```
In [32]: reduceLRonPlateau_cnn_reg_callback = tf.keras.callbacks.ReduceLRonPlateau(monitor = "val_1",
                                                                              patience = 3, verbose = 0)
```

```
In [33]: def get_CNN_reg_model(wd, rate, input_shape):
    model = Sequential([
        Conv2D(filters = 16, input_shape = input_shape, kernel_size = (3, 3), padding = 'SAME',
        MaxPooling2D(pool_size=(2, 2), name='pool_1'),
        BatchNormalization(),
        Conv2D(filters = 8, kernel_size = (3, 3), padding = 'SAME', kernel_regularizer = regularizers.l2_regularizer(wd),
        MaxPooling2D(pool_size=(2, 2), name='pool_2'),
        BatchNormalization(),
```

```

Conv2D(filters = 8, kernel_size = (3, 3), padding = 'SAME', kernel_regularizer = regul
Flatten(name='flatten'),
Dense(units=32, activation='relu', name='dense_1'),
Dropout(rate),
Dense(units=32, activation='relu', name='dense_2'),
Dense(units=10, activation='softmax', name='softmax')
])
return model

```

```
In [34]: CNN_reg_model = get_CNN_reg_model(1e-5, 0.4, train_x_std[0].shape)
```

```
In [35]: CNN_reg_model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv_1 (Conv2D)	(None, 32, 32, 16)	160
pool_1 (MaxPooling2D)	(None, 16, 16, 16)	0
batch_normalization (Batch Normalization)	(None, 16, 16, 16)	64
conv_2 (Conv2D)	(None, 16, 16, 8)	1160
pool_2 (MaxPooling2D)	(None, 8, 8, 8)	0
batch_normalization_1 (Batch Normalization)	(None, 8, 8, 8)	32
conv_3 (Conv2D)	(None, 8, 8, 8)	584
flatten (Flatten)	(None, 512)	0
dense_1 (Dense)	(None, 32)	16416
dropout (Dropout)	(None, 32)	0
dense_2 (Dense)	(None, 32)	1056
softmax (Dense)	(None, 10)	330
=====		

Total params: 19,802
Trainable params: 19,754
Non-trainable params: 48

```
In [36]: CNN_reg_model.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics=['acc
```

```
In [37]: CNN_reg_history = CNN_reg_model.fit(x=train_x_std, y=train_y,
epochs = 30,
validation_split=0.15, batch_size=64,
callbacks = [ checkpoint_val_reg_per_epoch, reduceLROnPlateau ],
verbose = False )
```

Epoch 00001: val_loss improved from inf to 0.97863, saving model to model_reg_checkpoints/checkpoint

Epoch 00002: val_loss improved from 0.97863 to 0.67881, saving model to model_reg_checkpoints/checkpoint

Epoch 00003: val_loss improved from 0.67881 to 0.64365, saving model to model_reg_checkpoints/checkpoint

nts/checkpoint

Epoch 00004: val_loss improved from 0.64365 to 0.55008, saving model to model_reg_checkpoints/checkpoint

Epoch 00005: val_loss did not improve from 0.55008

Epoch 00006: val_loss improved from 0.55008 to 0.51227, saving model to model_reg_checkpoints/checkpoint

Epoch 00007: val_loss improved from 0.51227 to 0.50802, saving model to model_reg_checkpoints/checkpoint

Epoch 00008: val_loss did not improve from 0.50802

Epoch 00009: val_loss improved from 0.50802 to 0.50006, saving model to model_reg_checkpoints/checkpoint

Epoch 00010: val_loss improved from 0.50006 to 0.46496, saving model to model_reg_checkpoints/checkpoint

Epoch 00011: val_loss did not improve from 0.46496

Epoch 00012: val_loss improved from 0.46496 to 0.46125, saving model to model_reg_checkpoints/checkpoint

Epoch 00013: val_loss did not improve from 0.46125

Epoch 00014: val_loss did not improve from 0.46125

Epoch 00015: val_loss did not improve from 0.46125

Epoch 00015: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.

Epoch 00016: val_loss improved from 0.46125 to 0.44587, saving model to model_reg_checkpoints/checkpoint

Epoch 00017: val_loss improved from 0.44587 to 0.42879, saving model to model_reg_checkpoints/checkpoint

Epoch 00018: val_loss did not improve from 0.42879

Epoch 00019: val_loss did not improve from 0.42879

Epoch 00020: val_loss did not improve from 0.42879

Epoch 00020: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.

Epoch 00021: val_loss improved from 0.42879 to 0.42211, saving model to model_reg_checkpoints/checkpoint

Epoch 00022: val_loss did not improve from 0.42211

Epoch 00023: val_loss did not improve from 0.42211

Epoch 00024: val_loss did not improve from 0.42211

Epoch 00024: ReduceLROnPlateau reducing learning rate to 8.000000525498762e-06.

Epoch 00025: val_loss improved from 0.42211 to 0.42208, saving model to model_reg_checkpoints/checkpoint

Epoch 00026: val_loss improved from 0.42208 to 0.42197, saving model to model_reg_checkpoints/checkpoint

Epoch 00027: val_loss improved from 0.42197 to 0.42132, saving model to model_reg_checkpoints/checkpoint

nts/checkpoint

Epoch 00028: val_loss did not improve from 0.42132

Epoch 00029: val_loss did not improve from 0.42132

Epoch 00030: val_loss improved from 0.42132 to 0.42129, saving model to model_reg_checkpoints/checkpoint

Epoch 00030: ReduceLROnPlateau reducing learning rate to 1.6000001778593287e-06.


```
In [38]: CNN_reg_model.evaluate(test_x_std, test_y, verbose = 2)
```

26032/1 - 7s - loss: 0.4076 - accuracy: 0.8637

```
Out[38]: [0.4697355834175343, 0.86370623]
```


```
In [39]: # Plot the training and validation loss

plt.plot(CNN_reg_history.history['loss'])
plt.plot(CNN_reg_history.history['val_loss'])
plt.title('Loss vs. epochs')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')
plt.show()
```




```
In [40]: plt.plot(CNN_reg_history.history['accuracy'])
plt.plot(CNN_reg_history.history['val_accuracy'])
plt.title('Accuracy vs Epochs')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(['Training', 'Validation'], loc = 'lower right')
```

```
Out[40]: <matplotlib.legend.Legend at 0x7fb67695e450>
```




```
In [41]: plt.imshow(test_x_std[3])
prediction = CNN_reg_model.predict(test_x_std[np.newaxis, 3])
print(np.argmax(prediction))
prediction
```

```
Out[41]: 0
array([[0.640227 , 0.02908207, 0.01007334, 0.01061409, 0.00300272,
        0.00406485, 0.1213654 , 0.0096774 , 0.10857343, 0.06331967]],
      dtype=float32)
```



```
In [42]: plt.imshow(test_x_std[769])
prediction = CNN_reg_model.predict(test_x_std[np.newaxis, 769])
print(np.argmax(prediction))
prediction
```

```
Out[42]: 3
array([[4.85487953e-06, 1.02035818e-03, 1.12710826e-04, 9.76971805e-01,
        6.33798263e-05, 2.08731536e-02, 3.09292023e-04, 2.14430620e-04,
        3.19963292e-04, 1.09997345e-04]], dtype=float32)
```




In [43]:

```
plt.imshow(test_x_std[2349])
prediction = CNN_reg_model.predict(test_x_std[np.newaxis, 2349])
print(np.argmax(prediction))
prediction
```

Out[43]:

```
0
array([[9.7721785e-01, 3.9058824e-03, 2.9240371e-04, 1.9991955e-04,
        2.4518691e-05, 6.3414141e-06, 2.2798381e-03, 4.6642619e-04,
        7.4521159e-03, 8.1548002e-03]], dtype=float32)
```



4. Get model predictions

- Load the best weights for the MLP and CNN models that you saved during the training run.
- Randomly select 5 images and corresponding labels from the test set and display the images with their labels.
- Alongside the image and label, show each model's predictive distribution as a bar chart, and the final model prediction given by the label with maximum probability.

In [44]:

```
print(checkpoint_path)
print(checkpoint_cnn_reg_path)
```

```
model_checkpoints/checkpoint
model_reg checkpoints/checkpoint
```

In [45]:

[illegible]

```
        metrics=['accuracy'])
reloaded_MLP_reg_model.evaluate(test_x_std, test_y, verbose = 2)
```

Out[45]: 26032/1 - 2s - loss: 2.3342 - accuracy: 0.1167
[2.3113791253864218, 0.116664104]

In [46]: *# Load Best MLP_Reg Model using the 'best' checkpoint we earlier saved*

```
reloaded_MLP_reg_model.load_weights(checkpoint_path)
reloaded_MLP_reg_model.evaluate(test_x_std, test_y, verbose = 2)
```

Out[46]: 26032/1 - 2s - loss: 0.9941 - accuracy: 0.7295
[0.884822539266224, 0.7295252]

In [47]: *# Create a new model and evaluate it without loading the previous checkpoints*

```
reloaded_CNN_reg_model = get_CNN_reg_model(1e-5, 0.4, train_x_std[0].shape)
reloaded_CNN_reg_model.compile(optimizer='adam',
                               loss='categorical_crossentropy',
                               metrics=['accuracy'])
reloaded_CNN_reg_model.evaluate(test_x_std, test_y, verbose = 2)
```

Out[47]: 26032/1 - 7s - loss: 2.3111 - accuracy: 0.0758
[2.306050506778194, 0.075829744]

In [48]: *# Load Best CNN_Reg Model using the 'best' checkpoint we earlier saved*

```
reloaded_CNN_reg_model.load_weights(checkpoint_cnn_reg_path)
reloaded_CNN_reg_model.evaluate(test_x_std, test_y, verbose = 2)
```

Out[48]: 26032/1 - 7s - loss: 0.4076 - accuracy: 0.8637
[0.4697355834175343, 0.86370623]

In [49]: *# For simplicity we will leave the prediction of 10 to correspond to 0th index*

```
labels = [
    '10',
    '1',
    '2',
    '3',
    '4',
    '5',
    '6',
    '7',
    '8',
    '9'
]
```

In [50]:

```
def show_table_of_n_random_data_labels_and_pred_dist(model, model_type, n):

    for i in range(0, n):

        random_index = np.random.choice(test_x_std.shape[0])
        prediction = model.predict(test_x_std[np.newaxis, random_index])
        ht = np.reshape(prediction, (10,))
        fig, (ax1, ax2) = plt.subplots(ncols = 2, nrows = 1, figsize = (10,1))
        ax1.imshow(test_x_std[random_index])
```



```

ax1.set_title('Actual Label: {}'.format(non_encoded_test_y[random_index]))
ax2.bar(range(0,10), height = ht)
ax2.set_title('Prediction: {} from {} Prediction Distribution'.format(np.argmax(h),
plt.show()

```


In [52]:

```


show_table_of_n_random_data_labels_and_pred_dist(MLP_reg_model, "MLP", 5)
show_table_of_n_random_data_labels_and_pred_dist(CNN_reg_model, "CNN", 5)

```


Actual Label: 4




Prediction: 4 from MLP Prediction Distribution




Actual Label: 0




Prediction: 0 from MLP Prediction Distribution




Actual Label: 8




Prediction: 8 from MLP Prediction Distribution




Actual Label: 3




Prediction: 3 from MLP Prediction Distribution




Actual Label: 1




Prediction: 1 from MLP Prediction Distribution




Actual Label: 1




Prediction: 1 from CNN Prediction Distribution




Actual Label: 4




Prediction: 9 from CNN Prediction Distribution




Actual Label: 3




Prediction: 5 from CNN Prediction Distribution




Actual Label: 4




Prediction: 4 from CNN Prediction Distribution



Actual Label: 9



Prediction: 9 from CNN Prediction Distribution



In []: